

# Brian Wheatman

[brianwheatman.com](http://brianwheatman.com)

## EDUCATION

### Johns Hopkins University (JHU)

PhD Computer Science; Advisor: Randal Burns

Dissertation: Cache and Memory Optimized Data Structures for High Performance Applications

Baltimore, MD

December 2024

### Massachusetts Institute of Technology (MIT)

Master of Engineering in Computer Science and Engineering

Concentration in Systems

Cambridge, MA

June 2019

### Massachusetts Institute of Technology (MIT)

Bachelor of Science in Computer Science and Engineering

Bachelor of Science in Mathematics - Minor in Economics - Minor in Management Science

Cambridge, MA

June 2017

## RESEARCH EXPERIENCE

### University of Chicago Department of Computer Science

Postdoctoral Research Scholar

Chicago, IL

September 2024 - Present

- Designing and implementing parallel algorithms for graph processing on proposed, next-generation, sparse supercomputers, capable of scaling to millions of cores

### JHU Department of Computer Science

PhD Researcher with Professor Randal Burns

Baltimore, MD

January 2020 - September 2024

- Designed and implemented parallel data structures to efficiently store dynamic graphs and matrices, which reduced the memory bandwidth requirements of analysis and computations run on these structures

### JHU Distributed Systems and Networks Lab

PhD Researcher with Professor Yair Amir

Baltimore, MD

September 2019 - September 2021

- Investigated methods for providing theoretical guarantees in artificial intelligence systems
- Designed a traffic light control algorithm, which outperforms static algorithms using machine learning without suffering from edge cases, as in standard machine learning approaches

### MIT Computer Science and Artificial Intelligence Laboratory (CSAIL)

Master of Engineering Thesis with Professor Charles Leiserson

Cambridge, MA

September 2017 - June 2019

- Improved an image processing pipeline containing petabytes of images with approximate locations to create a single, large mosaic with accurate locations
- Decreased time and resource requirements allowing the pipeline to run on one server instead of a cluster

### MIT Computer Science and Artificial Intelligence Laboratory (CSAIL)

SuperUROP Researcher with Professor Daniela Rus and Professor Sertac Karaman

Cambridge, MA

September 2016 - June 2017

- Designed and implemented a new, online algorithm for a NP-Hard variant of the Traveling Salesperson Problem, which enabled solutions ten times faster than previous approaches

### MIT Human Dynamics Laboratory at MIT Media Laboratory

Undergraduate Researcher with Professor Alex (Sandy) Pentland

Cambridge, MA

September 2015 - June 2016

- Modeled energy use from mobile phone data using machine learning
- Created predictive models for population movement using mobile phone records

### MIT Sloan School of Management

Undergraduate Researcher with Professor Stephen Graves

Cambridge, MA

February 2014 - August 2014

- Designed and developed a production line simulator to be used as a visualization and teaching tool

## TEACHING EXPERIENCE

Teaching Assistant/Course Assistant for Parallel Programming (JHU 601.320)

Fall 2021/Fall 2022/Fall 2023

Teaching Assistant for Intermediate Programming (JHU 601.220)

Fall 2020

Teaching Assistant for Computation Structures (MIT 6.004)

Spring 2018/Fall 2018/Spring 2019

Teaching Assistant for Seminar in Undergraduate Advanced Research (MIT 6.UAR)

Fall 2017

## PROFESSIONAL EXPERIENCE

### Lawrence Berkeley National Laboratory

*Visiting Researcher with Aydın Buluç*

Summer 2022

- Researched improvements to memory-efficient data structures, which improved their updatability and point query performance, while maintaining their near optimal scan performance
- Focused on exceptional, practical performance, while maintaining theoretical guarantees

### Google

*Software Engineering Intern*

Summer 2023

- Developed a simulator for a global network to investigate new strategies to increase reliability

*Software Engineering Intern*

Summer 2021

- Investigated distributed processing frameworks to select one for a distributed solver
- Implemented the distributed solver on the selected distributed processing framework

*Software Engineering Intern*

Summer 2020

- Investigated a hierarchical traffic engineering system to simplify routing
- Simulated the new system with past data to determine the impact of the hierarchical system

*Software Engineering Intern*

Summer 2019

- Designed a machine learning model for use in a proprietary truth inference problem
- Implemented the truth inference model and a distributed pipeline for preprocessing the data

*Software Engineering Intern*

Summer 2018

- Worked on Akaros, a new operating system for high performance and real time applications
- Developed a port of the Go programming language to run on Akaros

*Software Engineering Intern*

Summer 2017

- Enhanced a large, Map-Reduce data processing pipeline to aggregate and report advertising data
- Created a software cache for fast access to critical path metadata stored in slow storage

### Five Rings Capital

New York, NY

*Software Engineering Intern*

January 2017

- Conducted an evaluation of database solutions based on cost, scalability, and ease of use
- Reduced costs for a cloud database by structuring data to reduce the memory requirement

### Facebook

Menlo Park, CA

*Data Engineering Intern*

Summer 2016

- Created pipelines to collect and organize third party data for business intelligence purposes
- Worked with multiple large databases using several database engines, including Oracle and Hive
- Created self-documenting code that generated textual and visual aids to better understand the pipeline

### JP Morgan Chase & Co.

New York, NY

*Software Engineering Intern*

Summer 2015

- Designed and implemented an automated testing process for a large transaction system with nightly generated statistics and visualizations of performance
- Created a website to display results of the tests and long term trends

## AWARDS AND HONORS

- HPEC 2024 Outstanding Paper Award
- PPOPP 2024 Best Artifact Award
- HPEC 2020 Best Student Paper Award
- Johns Hopkins Gordon Croft Fellow (PhD Fellowship)
- Lockheed Martin Undergraduate Research and Innovation Scholar
- Member of Eta Kappa Nu (Computer Science Honor Society)
- Member of Tau Beta Pi (Engineering Honor Society)
- Boy Scouts of America Eagle Scout

## PUBLICATIONS

For more information see [brianwheatman.com](http://brianwheatman.com)

Brian Wheatman, Randal Burns, and Helen Xu, "Batch-Parallel Compressed Sparse Row: A Locality-Optimized Dynamic-Graph Representation" 2024 IEEE High Performance Extreme Computing Conference (HPEC). IEEE, 2024

Wheatman, Brian, et al. "BYO: A Unified Framework for Benchmarking Large-Scale Graph Containers." VLDB 2024

T. Kaler et al., "Speedcode: Software Performance Engineering Education via the Coding of Didactic Exercises," 2024 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW), 2024

Brian Wheatman, Randal Burns, Aydin Buluc, and Helen Xu, "CPMA: An Efficient Batch-Parallel Compressed Set Without Pointers" ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP) 2024

Helen Xu, Amanda Li, Brian Wheatman, Manoj Marneni, and Prashant Pandey, "BP-tree: Overcoming the Point-Range Operation Tradeoff for In-Memory B-trees" VLDB 2023

Brian Wheatman, Randal Burns, Aydin Buluc, and Helen Xu, "Optimizing Search Layouts in Packed Memory Arrays" 2023 SIAM Symposium on Algorithm Engineering and Experiments (ALENEX), 2023

Brian Wheatman and Randal Burns, "Streaming Sparse Graphs using Efficient Dynamic Sets" 2021 IEEE International Conference on Big Data (Big Data), 2021

Prashant Pandey, Brian Wheatman, Helen Xu, and Aydin Buluc, "Terrace: A Hierarchical Graph Container for Skewed Dynamic Graphs" Proceedings of the International Conference on Management of Data (SIGMOD), 2021

Brian Wheatman, et al. "RADICS: Runtime Assurance of Distributed Intelligent Control Systems" 2021 51st Annual IEEE/IFIP International Conference on Dependable Systems and Networks Workshops (DSN-W). IEEE, 2021.

Brian Wheatman and Helen Xu. "A Parallel Packed Memory Array to Store Dynamic Graphs" 2021 Workshop on Algorithm Engineering and Experiments (ALENEX). Society for Industrial and Applied Mathematics, 2021

Tim Kaler, Brian Wheatman, and Sarah Wooders. "High-Throughput Image Alignment for Connectomics using Frugal Snap Judgments" 2020 IEEE High Performance Extreme Computing Conference (HPEC). IEEE, 2020

Brian Wheatman and Helen Xu. "Packed compressed sparse row: A dynamic graph representation" 2018 IEEE High Performance Extreme Computing Conference (HPEC). IEEE, 2018

Brian Wheatman, Alejandro Noriega, and Alex Pentland. "Electricity demand and population dynamics prediction from mobile phone metadata" SBP-BRiMS 2016

## PRESENTATIONS

*Recent Improvements to Packed Memory Arrays*, Invited talk at Workshop on Recent Advances in Parallel and Concurrent Data Structures at SPAA 2024

*Ordered Sets: An Evolution of Memory Optimized Data Structures*, Invited talk at University of Chicago 2024

*So You Want to Make a Dynamic Graph Data Structure*, Invited talk at Workshop on Large-Scale Graph Processing at SPAA 2022

Guest Lecturer on SIMD and Vectorization (University of MD CMSC858N, Spring 2023; Georgia Tech CSE 6230, Spring 2024)

## COMMUNITY SERVICE

I have served on the following committees:

- Fastcode Programming Challenge (FCPC 2025) at PPoPP 2025
- Highlights of Parallel Computing (HOPC 2024) at SPAA 2024
- ACM SIGPLAN Principles and Practice of Parallel Programming (PPoPP 2024, 2025) Artifact Evaluation

I have served as a reviewer for the following conferences:

- ESA 2023, SEA 2022 and 2023, SPAA 2022 - 2024, SICOMP 2023, IPDPS 2023 and 2025, ALENEX 2023, BigData 2022, FOCS 2022, ACDA 2021

I contribute to the OpenCilk parallel programming platform.